

CSC 143 Java

Hashing

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-1

Review

- Want to implement Sets of objects
 - Want fast `contains()`, `add()`
- One strategy: a sorted list
 - OK `contains()`: use binary search
 - Slow `add()`: have to maintain list in sorted order
- Another strategy: a binary search tree
 - OK `contains()`: use binary search through tree
 - OK `add()`: use binary search to find right place to insert

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-2

A Magical Strategy

- What if... we had a magic method that could *convert each possible element value into its own unique integer*?
 - Takes an element, returns an integer (called a *hash code*)
 - Called a *perfect hash function*
- Then we could store the set elements in an array, with each element stored at an index equal to its hash code



- Array access is constant time – very fast: $O(1)$
- If computing the hash value is also $O(1)$, lookup is $O(1)$
 - Beats $O(\log n)$, which is the best we've seen so far

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-3

Hash Function Example

- Suppose we wanted to hash on a person's last name
- Use the individual characters of the name to compute a number
 - Example: cast each char to its int value, add all the int values
- Use the integer as an index into an array
- Drawbacks?
 - Array would be very large
 - "Soto" and "Soot" hash to the same value
Called a "collision"
- There are better string hash functions

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-4

If Only We Had A Perfect Hash...

- A *Perfect* hash function is one which has no collisions
 - two different objects never have the same hash code

How fast is `contains()`?

- Would just test whether value at the hash location index was non-null
- Fast!

How fast is `add()`?

- would just set the index to contain the element
- Fast!

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-5

Perfect vs. Imperfect Hash Functions

- *Perfect* hash functions are practical to implement only in limited cases
 - When the set of possible elements is small and known in advance
- But *imperfect* hash functions are practical
- An *imperfect (or regular) hash function* can produce collisions
- Imperfect hash functions compromise the promise of fast performance
 - How?
 - Can we salvage the design?

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-6

Solution: Buckets

- Instead of each array position containing the set elements directly...
 - it can contain a *list* of elements that all share the same hash code
 - This list is called a *bucket*
 - Unlike ordinary buckets, this kind can never be full!
- To test whether an element is in the set:
 - Use the hash code to find the correct bucket
 - Search that bucket's list for the element
- Add works similarly



3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-7

More about Buckets

- If hash function is good, then most elements will be in different buckets, and each bucket will be short
 - Most of the time, contains() and add() will be fast!
- There will probably be unused buckets – particularly at first
 - No data value happens to hash to a particular bucket
- Tradeoff:
 - more buckets: shorter linked lists, more unused space
 - fewer buckets: longer linked lists, less unused space
- Footnote: This approach is called *open hashing* (or *chaining*); there is another variation called *closed hashing* (or *open addressing*).

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-8

Object Hash Codes in Java

- Class Object defines a method hashCode() which returns an integer code for an object
- Strives to be different for different objects, but might not always be
 - Generally, you should assume the default hashCode in Java is very imperfect
- Subclasses can override this if a more suitable hash function is appropriate for instances

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-9

Hash Codes in Your Own Classes

- Subclasses should override hashCode() if a more suitable hash function is appropriate for instances
- Key rule: if o1 and o2 are different objects, then if


```
o1.equals(o2) == true
```

 it must also be true that


```
o1.hashCode() == o2.hashCode()
```
- Corollary: If you override either of hashCode() or equals(...) in a class, you probably should override the other one to be consistent
- **Danger:** The Java system cannot enforce these rules. A well-designed ("proper") class will follow them as a matter of good practice

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-10

HashCode for Complex Objects in Java

- Key idea: calculate a hash code value using the fields that are considered in method equals
- Hash codes for individual fields
 - Boolean: 0 or 1; int, char: cast to int; float, double, long: get the bits (see ref.)
 - Object reference: assuming this field implements equals by recursively calling equals on its parts, call get the hashCode for the fields
- Combining the field hash codes – one possibility


```
result = 17;
for each hash code c for some part of the object, set result = 37*result+c;
return result
```
- Source: *Effective Java* by Joshua Bloch (A-W, 2001)

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-11

HashMap: Java Library Dictionary Class

- The java.util.HashMap implements a dictionary using a hash table
 - Uses the objects hashCode() method
- Key operations (interface Map)

```
public interface Map<K,V> {
    // Associate the given key with the given value (return previous value or null)
    public V put(K key, V value);
    // Return the value associated with the key, or null if no such value
    public V get(Object key);
    // Remove the key and its associated object from the map
    public V remove(Object key);
}
```

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-12

Implementing a HashSet Class

• HashSet: an implementation of Set using hashing

```
public class HashSet implements Set {
    private List[] buckets; // buckets[k] is a list of elements that satisfy
                           // elem.hashCode() % nBuckets == k
                           // buckets[k]==null if no elems have hashcode k

    private static final int nBuckets = 101; // default # of buckets
    public HashSet() {
        buckets = new List[nBuckets]; // each elem initialized to null
    }
    ...
}
```

- Generally, having a prime number of buckets produces a decent distribution of objects among the buckets

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-13

Computing the Bucket Number

• Algorithm:

- Compute the object's hash code
- Convert it into a legal index into the buckets array: something in the range 0..buckets.length-1

```
/** Return the index in buckets where the elem would be found, if it's in the set */
private int bucketNum(Object elem) {
    return elem.hashCode() % buckets.length;
}
```

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-14

Adding a New Element

```
public boolean add(Object elem) {
    int i = bucketNum(elem);
    List bucket = buckets[i];
    if (bucket == null) {
        // this is the first element in this bucket; create the bucket list first
        bucket = new ArrayList();
        buckets[i] = bucket;
    } else { // return false if elem is already contained in the set
        if (bucket.contains(elem)) { return false; }
    } // otherwise add element to bucket's list
    bucket.add(elem);
    return true;
}
```

- Note that this (and following) code relies on fact that array elements are null when an array is first created

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-15

Checking Whether an Element is In the Set

```
public boolean contains(Object elem) {
    int i = bucketNum(elem);
    List bucket = buckets[i];
    if (bucket == null) {
        // empty bucket
        return false;
    } else {
        // look for element in non-empty bucket
        return bucket.contains(elem);
    }
}
```

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-16

How Efficient is HashSet?

• Parameters

- n number of items stored in the HashSet
- b number of buckets

• Load factor: n/b – ratio of # entries to # buckets

• Cost of contains() and add() is roughly constant, independent of the size of the set, provided that:

- Hash function is good – distributes keys evenly throughout buckets
Ensures that buckets are all about the same size; no really long buckets
- Load factor is small
Don't have to search too far in any bucket

• In the average case, the fastest set implementation!

- In the worst case, the slowest...

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-17

Some Issues

• Interesting issues for data structures courses

- How do you pick a good hash function?
Needs to be $O(1)$ and produce few duplicates
- How do you keep the load factor small?
One answer: Grow the buckets array and rehash all the elements if the load factor gets too large

- Take a junior-level data structures course (e.g. CSC 273 at NSCC, or maybe CSE373 or CSE326 at UW Seattle) to learn more!

3/14/2011

Original © University of Washington, used by permission, modified by Vince Offenback

22-18

Summary

- Hash functions "guess" the right index to look for an element
 - Can do it faster than binary search can
- If most buckets are short (e.g. ≤ 3 elements), then works very well
- To keep buckets small, need:
 - good hash functions and
 - the ability to grow the buckets array

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-19

Comparing Data Structures

- We now have several implementations of data structures in which we can store and search for objects
 - Array-based lists
 - Linked lists
 - Trees
 - Binary search trees, in particular
 - Hash sets
- Each offers various tradeoffs of performance for common operations
 - Add, remove, contains, iterate (either in random or sequential order)
- Which one is best? It depends on the application!

3/14/2011

Original © University of Washington; used by permission; modified by Vince Offenback

22-20